

5 Steps to Speed Up Your Data-Analysis on a Single Core


Jonathan Striebel

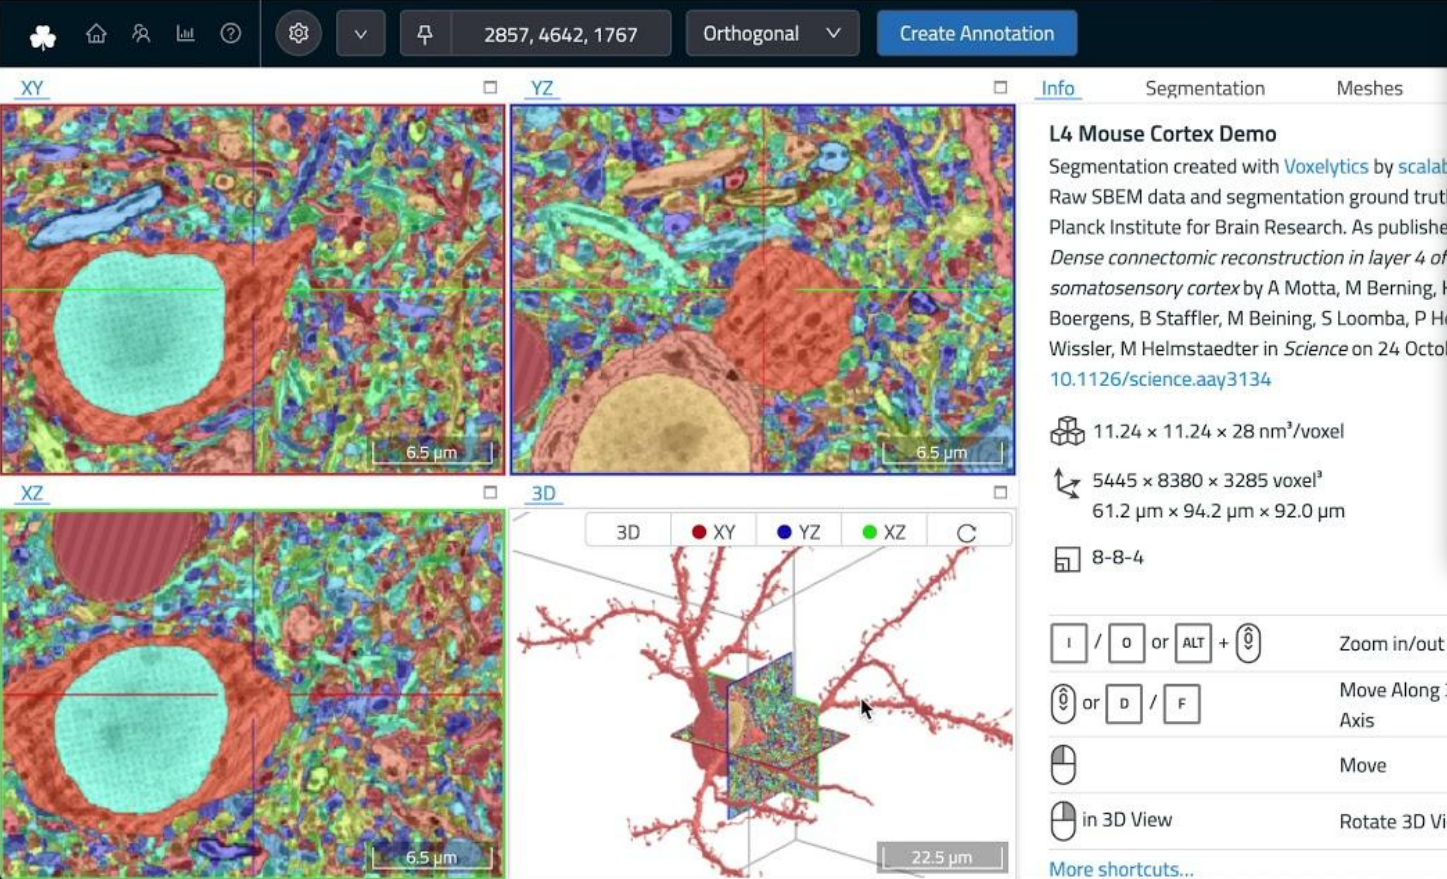


Hi, I'm **Jonathan Striebel**.

 @jostriebe

jonathan@scalableminds.com






Python library: `pip install webknossos`
see docs.webknossos.org/webknossos-py

Large Scale Data-Analysis Experiments



PB-scale
image data

Machine Learning Systems




Segmentation & Agglomeration
Algorithms


Running weeks in
HPC Clusters



Neuron
Reconstructions for
Biological Analysis




Why to Speed Up



It's too slow

Why to Speed Up: Example






Combining statistics of billions of segments across thousands of chunks

1. ∞ (crashes with OOM after 2 days, 1TB memory)
2. 18h (3-4 iterations per week)
3. 7h (5-8 iterations per week)

500+% faster IRL

Toy Example: ~200% faster



Why to Speed Up on a Single Core

Don't parallelize (yet):

- Single core improvements **pay off, also when parallelized** later
- Parallelization **needs resources**
(cores, memory, cluster-nodes, money, time)
- Code may be **hard / impossible** to parallelize
(e.g. when reducing results from a map-reduce)

5 Steps



Profiling



Efficient IO



Vectorization



Memory & Precision Tradeoffs



Jit-ting with numba



1. Profiling



Speed

- **py-spy** (sampling based, flamegraphs, speedscores)
- **yappy** (line-wise)
- **cProfile** (ships with Python)
- **pyinstrument**
- **Palanteer**

Memory

- **memory-profiler**
- **Guppy3**


2. Efficient IO

Text-based

csv

json

yaml



Binary format

hdf5

npz


parquet

pickle

sqlite

zarr

...



50-70% faster

see also


github.com/mverleg/array_storage_benchmark

3. Vectorization



NumPy

compact representation for numerical arrays
with optimized functions



170-180% faster

Pandas



tabular data backed by numpy

often semi-fast, consider alternatives:

Polars, datatable, PandasPy, modin (multi-threaded), vaex (out-of-core), dask-dataframe (parallel)

4. Memory & Precision Tradeoffs

- **Data type**
 - direct effect on memory & precision
 - time affected via read/write/operations on data
- **Iterative methods (e.g. divide & conquer, reduce)**
 - less data to keep in memory
 - iterations might lose numerical precision
- **Lookup tables**
 - time vs memory
- **Compression**
 - lossless: time vs memory
 - lossy: time vs memory vs precision




180-200% faster

5. Jit-ting with numba

1. Add `@numba.njit` decorator, enables jitting in nopython mode

2. Fix shape & dtype errors, as broadcasting is more strict in numba

3. 🚀



190-200% faster

Further Steps

- Upgrade dependencies & Python
- Faster Python Runtimes
PyPy, Pyjion, Cinder
- Optimize critical code-paths closer to the metal
Cython, pybind11, cffi, PyO3, ONNX, ...
- Parallelization
async, threading, multiprocessing, Spark/Dask/Ray, ...

Summary



Profiling



Efficient IO



Vectorization



Memory & Precision Tradeoffs



Jit-ting with numba



Jonathan Striebel | jonathan@scalableminds.com

@jostriebe

Thanks!

code @

github.com/jstriebel/data-analysis-speedup

Interested in working at scalable minds?

We're looking for Python Engineers, Scala Backend Engineers

Jonathan Striebel | jonathan@scalableminds.com @jostriebe

